

CyberSentinel: AI-Driven Real-Time Cybersecurity Solution

Vaishnavi Dasari^{*}, Snehitha Bodiga[†], Shoaib Khan Patan[‡], Abisainath Arepalli[§],
Aasritha Bhimisetty[¶], Neelkamal Rana^{||}, Dhanush Kolapalli^{**}

Department of Computer Science, Pace University, New York, USA

Emails: ^{*}vd73531n@pace.edu, [†]sb65320n@pace.edu, [‡]sp05764n@pace.edu,

[§]aa98295n@pace.edu, [¶]ab74717n@pace.edu, ^{||}nr37172n@pace.edu, ^{**}dk25992n@pace.edu

Abstract—CyberSentinel is an advanced cybersecurity platform built using artificial intelligence and machine learning techniques to provide real-time detection and prevention of digital threats. It focuses on defending individuals and enterprises against phishing attacks, identity theft, fraudulent transactions, and malicious web links. The system integrates behavioral analytics, NLP-based email analysis, anomaly detection, and cloud-based threat intelligence to offer dynamic, adaptive security. Designed with scalability and accessibility in mind, CyberSentinel ensures that even non-technical users can benefit from intelligent cybersecurity defenses.

Index Terms—Cybersecurity, Artificial Intelligence, Machine Learning, Phishing Detection, Fraud Detection, NLP, Threat Intelligence

I. INTRODUCTION

In an increasingly digital world, cybersecurity threats are growing in complexity and frequency. With the rise of cloud computing, remote work, and interconnected systems, the attack surface has expanded exponentially. Cybercriminals are leveraging sophisticated tools and techniques, making it harder for traditional rule-based security systems to keep up. These legacy systems, while effective in handling known threats, often fail to adapt to evolving cyber attack vectors.

Traditional cybersecurity approaches rely heavily on static signatures and predefined rules to identify malicious behavior. While these methods are computationally efficient, they lack the adaptability required to detect novel attacks such as spear-phishing, zero-day vulnerabilities, and polymorphic malware. Furthermore, they often generate high false-positive rates, leading to alert fatigue among analysts and missed real threats.

To address these limitations, CyberSentinel introduces a dynamic, AI-driven security platform that emphasizes proactive defense. By utilizing machine learning (ML) and natural language processing (NLP), the system continuously learns from new data and user feedback, making it capable of detecting anomalies and previously unseen attack patterns in real time. This intelligent design enhances the ability to provide early warnings, precise threat classifications, and actionable insights.

The platform is built to serve a wide range of users—from everyday email recipients to enterprise-level security administrators. CyberSentinel offers a unified dashboard where users can upload suspicious emails, check the safety of URLs, and monitor transactions for fraud. With its user-friendly interface and scalable cloud-based backend, the platform empowers

users to take control of their digital security with minimal technical expertise.

II. LITERATURE REVIEW

Cybersecurity research over the last two decades has extensively explored the limitations of traditional detection systems. Signature-based detection mechanisms, which depend on known attack patterns, are commonly used in antivirus software and intrusion detection systems (IDS). However, these systems struggle against zero-day threats—attacks that exploit previously unknown vulnerabilities. They also require constant manual updates, making them impractical for fast-evolving cyber landscapes.

Recent advancements in artificial intelligence have introduced more robust approaches to threat detection. Natural language processing (NLP), for instance, has been effectively applied to phishing email detection by analyzing patterns, tone, and intent within message content. Techniques such as tokenization, embedding, and transformer-based models (e.g., BERT) have significantly improved classification accuracy and contextual understanding, making NLP a cornerstone in modern email security systems.

Similarly, anomaly detection techniques such as Autoencoders and Isolation Forests have been employed in the detection of fraudulent financial transactions. Unlike rule-based methods, these models learn the normal behavior of users or systems and flag deviations that may indicate fraud or unauthorized access. This unsupervised learning capability is crucial for identifying new fraud schemes that may not have been seen before.

CyberSentinel builds upon these AI-driven techniques by combining NLP for phishing detection, anomaly detection for financial fraud, and deep learning models for malicious URL identification. The integration of these models into a unified platform offers a holistic solution that not only detects threats with high accuracy but also adapts over time. By leveraging real-time data and user feedback, CyberSentinel ensures continuous learning and improvement, addressing key challenges highlighted in existing literature.

III. SYSTEM DESIGN AND ARCHITECTURE

CyberSentinel is built using a modular, service-oriented architecture that enables high scalability, maintainability, and

extensibility. The system is designed with a clear separation of concerns, allowing each component to operate independently while maintaining smooth integration through a well-defined API structure. This architecture supports both microservice deployment and monolithic development environments, depending on the needs of the host infrastructure.

The frontend is developed using React.js to provide a modern, responsive user interface. It offers intuitive workflows for users to upload and classify emails, submit transactions for fraud detection, and verify URLs. The UI also features dashboards for monitoring system statistics, managing users and roles, and responding to alerts. This client-side application communicates with the backend over RESTful APIs secured with OAuth 2.0 tokens.

The backend, developed with Flask (Python), handles data processing, model inference, user management, and API interactions. It integrates with a PostgreSQL database to store user input, classification results, system logs, and access credentials. Each machine learning model—phishing classifier, fraud detector, and URL analyzer—is containerized for isolated deployment and connected through internal API gateways.

Cloud deployment is managed via Amazon Web Services (AWS), leveraging services like EC2 for compute, S3 for storage, and RDS for database management. The cloud infrastructure ensures horizontal scalability, allowing CyberSentinel to handle concurrent users and data submissions efficiently. OAuth 2.0 serves as the backbone of authentication and authorization, ensuring role-based access control and session management across all system layers.

IV. METHODOLOGY

The development of CyberSentinel followed the Agile software development methodology, which facilitated iterative progress, flexibility in incorporating stakeholder feedback, and rapid prototyping. The team operated in three structured sprints, each lasting approximately four weeks. Each sprint ended with a demonstration and retrospective session, ensuring alignment with user expectations and performance objectives.

A. Agile Sprint Structure

Sprint 1: Foundational Design and Setup

The initial phase focused on setting up core infrastructure and planning. Key deliverables included:

- Designing UI wireframes using Figma to outline user journeys.
- Setting up Git repositories and CI/CD pipelines using GitHub Actions.
- Deploying PostgreSQL and Flask on local Docker containers.
- Defining APIs and model integration strategies.

Sprint 2: Model Integration and Backend API Development

This sprint focused on integrating machine learning models into modular APIs:

- Trained BERT, Autoencoder, and CNN-LSTM models on labeled datasets.

- Packaged models as microservices using Flask and Docker.
- Developed RESTful endpoints for model inference and validation.
- Connected the frontend to backend services using Axios for secure communications.

Sprint 3: Real-Time Alerts and Performance Optimization

The final sprint aimed to refine the system's reliability and responsiveness:

- Implemented role-based access using OAuth 2.0 and JWTs.
- Developed an alert engine for immediate phishing/fraud notifications.
- Integrated logging and audit trails for traceability.
- Conducted user acceptance testing (UAT) with test participants.

B. Model Development and Training

AI model development was central to the system's effectiveness. The development process included dataset curation, preprocessing, model selection, training, and evaluation.

Toolkits Used:

- **Scikit-learn** was used for implementing Isolation Forests and handling cross-validation, grid search, and preprocessing.
- **PyTorch** provided the flexibility and GPU acceleration necessary for training deep learning models like BERT and CNN-LSTM.

Model Pipelines:

- **Phishing Email Classifier:** Fine-tuned BERT on a corpus of phishing and legitimate emails. Tokenization, attention masking, and transfer learning were applied using Hugging Face's Transformers library.
- **Fraud Detection:** An unsupervised Autoencoder learned patterns in normal transaction data, while Isolation Forests were used to flag anomalies deviating from the learned distribution.
- **URL Classification:** A hybrid CNN-LSTM model was built. CNN layers captured spatial token patterns in URLs, and LSTM layers modeled sequential dependencies. This model was particularly effective at detecting obfuscated malicious links.

Evaluation Metrics: Each model was evaluated using:

- *Precision, Recall, F1-score* — to balance false positives and false negatives.
- *ROC-AUC* for the fraud detection model due to its imbalanced class distribution.
- *Confusion matrices and classification reports* during UAT to track misclassifications.

C. Security and Access Control

CyberSentinel integrates OAuth 2.0 for secure, token-based user authentication. User roles (admin, reviewer, general) are enforced through middleware policies and stored in the

PostgreSQL backend. This role-based access control (RBAC) ensures only authorized users access sensitive features like audit logs, retraining triggers, and flagged sample review.

D. Deployment and DevOps

Deployment emphasized scalability, maintainability, and security. Key practices include:

- **Containerization:** All services were containerized using Docker, ensuring consistency across development, testing, and production.
- **Cloud Hosting:** Services were deployed on AWS EC2 instances, while S3 handled static file storage and RDS supported relational data.
- **CI/CD Pipeline:** GitHub Actions was configured for automatic building, testing, and deployment to staging environments after code merges.
- **Monitoring:** Prometheus and Grafana were used to monitor system performance, model latency, and API uptime.

E. Secure Logging and Audit Trails

System logs, including login attempts, model predictions, user actions, and feedback events, were recorded and timestamped. Audit trails are stored securely and can be retrieved by authorized users for post-incident analysis. This feature not only supports compliance but also strengthens system integrity and accountability.

F. User-Centric Iteration and Feedback

Feedback was continuously incorporated from test users and advisors. Features such as real-time alerts, confidence score displays, and feedback buttons were developed based on usability testing. Logs of user corrections were used to fine-tune thresholds and retrain detection models in later phases.

V. IMPLEMENTATION

A. Modular Microservices

CyberSentinel is developed using a microservices-based architecture to promote scalability, flexibility, and fault isolation. Each service is independently developed, tested, deployed, and scaled, allowing for parallel development and quick fault recovery. The core services include:

- **Phishing Email Classifier:** Handles preprocessing of email content and classifies it using a fine-tuned BERT model.
- **Transaction Fraud Detector:** Accepts transaction data and evaluates it using an Autoencoder and Isolation Forest ensemble to detect anomalies.
- **URL Analyzer:** Performs string feature extraction on URLs and passes them through a CNN-LSTM model to identify potential malicious intent.
- **User Management System:** Manages user registration, authentication, role assignment, and session tracking via OAuth 2.0.
- **Notification Engine:** Sends alerts via email and updates the frontend with real-time threat notifications using WebSocket or REST callbacks.

Each microservice is containerized using Docker and orchestrated using Kubernetes for load balancing, health checks, and autoscaling based on usage.

B. Frontend Interface

The frontend is a single-page application built using React.js and Material UI for design consistency and accessibility. It is designed for both technical and non-technical users, with an emphasis on clarity and usability. Key features include:

- Drag-and-drop email or URL submission tools.
- Real-time feedback display with classification results.
- Graphical dashboards to monitor system activity and alert trends.
- Admin panel for managing users, model thresholds, and logs.

The interface uses JWTs obtained via OAuth 2.0 authentication to securely communicate with backend services through RESTful API endpoints.

C. Backend Services

The backend, written in Python using Flask, serves as the central orchestrator for business logic, validation, and model inference requests. It is structured into the following layers:

- **API Gateway:** Exposes RESTful endpoints, enforces rate limiting and request validation.
- **Service Handlers:** Route incoming requests to the appropriate microservice container for processing.
- **Persistence Layer:** Interfaces with a PostgreSQL database for storing user data, feedback, logs, and model outcomes.

Background jobs and scheduled tasks are handled using Celery and RabbitMQ, which offload compute-intensive retraining or batch scoring operations.

D. Model Integration

The platform integrates three main machine learning models, each tailored for a specific detection task:

- **Email Phishing Detection:** A fine-tuned BERT model trained on a labeled dataset of phishing and legitimate emails, with attention to context and semantics.
- **Fraudulent Transaction Detection:** An Autoencoder learns normal transaction behavior while Isolation Forest flags deviations. This hybrid approach increases robustness in imbalanced datasets.
- **Malicious URL Detection:** A CNN extracts spatial n-gram patterns, and the LSTM layer captures sequential dependencies, allowing the model to distinguish obfuscated malicious links from benign ones.

Models are packaged in isolated Docker containers, exposed as APIs, and monitored via Prometheus and Grafana for inference performance.

E. Role-Based Access Control

CyberSentinel employs a granular, role-based access control (RBAC) model secured via OAuth 2.0 and JWT. User roles include:

- **Admin:** Full access to all modules, including user management and model settings.
- **Reviewer:** Authorized to verify, approve, or provide feedback on flagged inputs.
- **General User:** Can submit data for classification and view personal history or system notifications.

All API endpoints are protected with role-aware middleware. Audit logs are generated for sensitive actions to support compliance and traceability.

F. Feedback and Notification System

To ensure model adaptability and user trust, CyberSentinel incorporates an interactive feedback mechanism:

- **Threat Alerts:** Users are instantly notified via email and dashboard popups when a submission is classified as malicious.
- **Feedback Loop:** Users can label false positives or negatives, which are stored in a retraining queue.
- **Model Retraining:** Periodically, new feedback is used to fine-tune the models, especially threshold-sensitive modules like Isolation Forests.
- **Dashboard Alerts:** Users and admins can view alert history and classification trends over time, allowing insights into threat evolution.

This system ensures that CyberSentinel improves over time and aligns with real-world patterns, reducing static decision-making limitations.

VI. TESTING AND EVALUATION

VII. TESTING AND EVALUATION

A. Testing Methodology

To ensure reliability and robustness, CyberSentinel underwent a multi-tiered testing process consisting of unit testing, integration testing, system testing, and user acceptance testing (UAT). The testing methodology was aligned with Agile sprint goals to validate functional correctness and non-functional requirements (e.g., performance, security, and usability).

Unit Testing: Each backend microservice and frontend component was tested in isolation using Pytest (for Python modules) and Jest (for React components). These tests focused on verifying logic accuracy, correct input validation, and appropriate exception handling.

Integration Testing: Once individual components passed unit testing, integration tests validated end-to-end workflows such as:

- Email upload → Phishing detection → Notification.
- Transaction entry → Fraud score retrieval.
- Role-based login → Dashboard routing and access controls.

Postman and Selenium were used to simulate frontend-to-backend interactions and verify API endpoint reliability.

User Acceptance Testing (UAT): A group of 15 diverse users including students, faculty, and IT professionals were involved in structured testing sessions. They evaluated real-time threat detection, system response, ease of use, and interpretability of outputs. Feedback was documented through questionnaires and recorded sessions.

Automation and CI Testing: Automated pipelines triggered test suites for each pull request using GitHub Actions. These scripts validated model predictions, database I/O, and token-based access control enforcement to ensure consistent behavior during continuous integration.

B. Model Performance

The performance of each AI model was evaluated against its respective task using both internal validation and unseen test datasets. Performance metrics included accuracy, precision, recall, F1-score, and ROC-AUC.

Phishing Email Detection: The fine-tuned BERT model achieved:

- **Accuracy:** 92.3%
- **Precision:** 91.7%
- **Recall:** 93.1%
- **F1-score:** 92.4%

The model demonstrated resilience to adversarial obfuscation techniques like link masking and subtle grammar alterations, making it highly practical for deployment.

Transaction Fraud Detection: The Autoencoder + Isolation Forest ensemble was tested on a synthetic but behaviorally realistic dataset. It achieved:

- **Anomaly Detection Accuracy:** 95.8%
- **False Positive Rate:** 2.4%
- **ROC-AUC:** 0.964

The model was particularly effective at identifying low-frequency fraud patterns without overfitting on high-frequency legitimate transactions.

URL Classification: The CNN-LSTM model evaluated URLs with embedded features such as token entropy and character distribution. It reached:

- **Accuracy:** 85.1%
- **Precision:** 84.0%
- **Recall:** 86.3%
- **Latency:** 420ms per prediction (avg)

Performance was evaluated on a benchmark dataset combining PhishTank and Alexa-ranked URLs.

C. System Evaluation

System-level evaluation focused on measuring performance under load, latency, scalability, and fault tolerance.

Latency and Throughput: The system handled an average of 100 concurrent requests per second without noticeable degradation in response time, supported by AWS EC2 autoscaling groups.

Scalability: Load tests using Apache JMeter confirmed linear scalability under increased input volume, primarily due to container-based deployment and stateless API design.

Fault Recovery: Chaos testing simulated node failures and verified that Kubernetes automatically re-routed traffic to healthy pods, maintaining 99.9% uptime over a 48-hour stress period.

Security Evaluation: Penetration tests were performed using OWASP ZAP and manually crafted token hijack attempts. All endpoints enforced proper access control, and token expiry mechanisms worked as expected.

D. User Feedback and Continuous Improvement

User feedback was collected throughout the testing phase to refine the system’s usability and relevance. Key observations included:

- **Dashboard Clarity:** Early test users found threat classifications helpful but requested clearer explanations. This led to the addition of color-coded severity levels and confidence scores.
- **Feedback Mechanism:** Users appreciated the ability to flag misclassifications. A retraining queue was implemented to aggregate this data for future model updates.
- **Alert System Enhancements:** Alerts were redesigned to include links to threat reports and risk scores, increasing transparency and response effectiveness.

Feedback data was stored in a PostgreSQL feedback table and linked with sample UUIDs. This dataset was later used in Sprint 3 to fine-tune model thresholds and improve sensitivity to edge cases.

The feedback loop thus forms a critical part of CyberSentinel’s adaptive learning system, ensuring it evolves alongside emerging cyber threat patterns and user behavior.

VIII. RESULTS AND DISCUSSION

The deployment of CyberSentinel highlighted the platform’s effectiveness in integrating artificial intelligence to detect and mitigate cyber threats. Across multiple testing stages and user scenarios, the models demonstrated strong generalization capabilities and maintained high accuracy even with varied and complex input data. The use of real-world phishing emails, synthetic fraud transaction datasets, and diverse URL samples ensured that the models were tested under realistic conditions.

The phishing detection module, powered by BERT, was able to accurately classify most phishing attempts, even when they were crafted with deceptive language or disguised links. Its contextual understanding significantly outperformed traditional keyword-based filters. The transaction fraud module showed particular strength in identifying subtle anomalies that might evade rule-based systems, highlighting the advantages of unsupervised learning in financial security.

The platform’s ability to adapt based on user feedback is another major achievement. The inclusion of a feedback loop enabled continuous learning, where flagged false positives and false negatives were used to refine model thresholds and retrain underperforming components. This adaptive behavior ensures that CyberSentinel evolves over time and reduces reliance on static configurations.

Challenges encountered during implementation included handling edge cases in text classification where legitimate emails shared traits with phishing attempts, leading to occasional misclassification. Additionally, ensuring the responsiveness of the API under concurrent requests required optimization at the container orchestration level. These issues were mitigated through iterative testing and improvements in model tuning and backend scaling strategies, ensuring consistent performance across a range of workloads.

IX. CONCLUSION AND FUTURE WORK

CyberSentinel successfully demonstrates how artificial intelligence can revolutionize modern cybersecurity practices by offering real-time, scalable, and intelligent threat detection. Its modular architecture, AI-powered classification models, and secure cloud deployment make it a versatile solution adaptable to different user needs—from individuals concerned about phishing emails to organizations handling sensitive financial data.

The project addressed core cybersecurity problems using state-of-the-art techniques in natural language processing, anomaly detection, and deep learning. With a user-friendly dashboard and real-time feedback mechanisms, CyberSentinel not only detects threats but empowers users to make informed decisions. The integration of role-based access and secure authentication further enhances its readiness for real-world deployment.

Looking ahead, several avenues for improvement have been identified. One key area is multilingual phishing detection, which would extend the platform’s usability across different regions and user bases. Expanding the phishing email model’s training dataset to include multiple languages and regional attack formats would significantly improve coverage.

Additional goals include integration with enterprise-level Security Operations Center (SOC) platforms to provide centralized threat intelligence, implementing automated model retraining pipelines, and optimizing the platform for mobile devices. These enhancements will further solidify CyberSentinel’s role as a comprehensive, next-generation cybersecurity tool that evolves with emerging threats and user expectations.

REFERENCES

- [1] Scikit-learn Documentation. Available: <https://scikit-learn.org>
- [2] PyTorch Tutorials. Available: <https://pytorch.org/tutorials>
- [3] Devlin, J., et al., “BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding,” arXiv preprint arXiv:1810.04805, 2018.
- [4] Kingma, D. P., and Welling, M., “Auto-Encoding Variational Bayes,” arXiv preprint arXiv:1312.6114, 2013.
- [5] Hochreiter, S., and Schmidhuber, J., “Long Short-Term Memory,” Neural Computation, vol. 9, no. 8, pp. 1735–1780, 1997.
- [6] Saxe, J., and Berlin, K., “Expose: A Character-Level Convolutional Neural Network with Embeddings for Detecting Malicious URLs,” arXiv:1702.08568, 2017.
- [7] OAuth 2.0 Protocol Specification. Available: <https://oauth.net/2/>
- [8] Pandas Documentation. Available: <https://pandas.pydata.org>
- [9] NumPy Documentation. Available: <https://numpy.org>
- [10] Matplotlib Documentation. Available: <https://matplotlib.org>
- [11] Seaborn Documentation. Available: <https://seaborn.pydata.org>
- [12] TensorBoard for Model Evaluation. Available: <https://www.tensorflow.org/tensorboard>

- [13] XGBoost Documentation. Available: <https://xgboost.readthedocs.io>
- [14] LightGBM Documentation. Available: <https://lightgbm.readthedocs.io>
- [15] SHAP for Explainable ML. Available: <https://github.com/slundberg/shap>
- [16] Lime for Model Interpretability. Available: <https://github.com/marcotcr/lime>
- [17] OWASP Phishing and Social Engineering. Available: <https://owasp.org>
- [18] PhishTank: Free Community Phishing Data. Available: <https://www.phishtank.com>
- [19] Google Safe Browsing API. Available: <https://developers.google.com/safe-browsing>
- [20] MITRE ATT&CK Framework. Available: <https://attack.mitre.org>
- [21] Lockheed Martin Cyber Kill Chain. Available: <https://www.lockheedmartin.com>
- [22] ENISA Threat Landscape 2023. Available: <https://www.enisa.europa.eu>
- [23] IBM X-Force Threat Intelligence Index 2023. Available: <https://www.ibm.com/reports/threat-intelligence>
- [24] Symantec Internet Security Threat Report, Volume 25, 2020.
- [25] Microsoft MSTIC Blog. Available: <https://www.microsoft.com/security/blog>
- [26] Palo Alto Unit 42 Threat Intelligence. Available: <https://unit42.paloaltonetworks.com>
- [27] Dua, D., and Graff, C., "UCI Machine Learning Repository," 2019. Available: <https://archive.ics.uci.edu/ml>
- [28] LeCun, Y., Bengio, Y., and Hinton, G., "Deep learning," *Nature*, vol. 521, pp. 436–444, 2015.
- [29] Goodfellow, I., Bengio, Y., and Courville, A., "Deep Learning," MIT Press, 2016.
- [30] Vaswani, A., et al., "Attention is All You Need," *NeurIPS*, 2017.
- [31] Brownlee, J., "Imbalanced Classification with Python," *Machine Learning Mastery*, 2020.
- [32] Tidymverse for R. Available: <https://www.tidyverse.org>
- [33] Ribeiro, M. T., et al., "Why Should I Trust You?: Explaining the Predictions of Any Classifier," *KDD*, 2016.
- [34] Chio, C., and Freeman, D., "Machine Learning and Security," *O'Reilly Media*, 2018.
- [35] Kim, Y., "Convolutional Neural Networks for Sentence Classification," *EMNLP*, 2014.
- [36] Zhang, Y., and Wallace, B. C., "A Sensitivity Analysis of CNNs for Sentence Classification," *EMNLP*, 2015.
- [37] Abadi, M., et al., "TensorFlow: A System for Large-Scale Machine Learning," *OSDI*, 2016.
- [38] Ross Anderson, "Security Engineering," 3rd ed., Wiley, 2020.
- [39] Kevin Mitnick, "The Art of Deception," Wiley, 2002.
- [40] OpenCV Documentation. Available: <https://opencv.org>
- [41] Meta AI Research Blog. Available: <https://ai.facebook.com/research>
- [42] Google AI Blog. Available: <https://ai.googleblog.com>
- [43] Shodan: Search Engine for Devices. Available: <https://www.shodan.io>
- [44] VirusTotal API. Available: <https://www.virustotal.com>
- [45] ISO/IEC 27001: Information Security Management, ISO, 2013.
- [46] Ponemon Institute, "Cost of a Data Breach Report 2023."
- [47] Jiang, M., et al., "Detecting Fraudulent Online Activities Using Graph Analysis," *ACM TOIT*, 2016.
- [48] Breiman, L., "Random Forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.
- [49] Mitrokovtsa, A., et al., "Intrusion Detection with Machine Learning Techniques," *Computers & Security*, 2012.
- [50] Ramachandran, P., et al., "Universal Sentence Encoder," *arXiv:1803.11175*, 2018.